Analyze_ab_test_results_notebook

September 28, 2020

0.1 Analyze A/B Test Results

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project RUBRIC. **Please save regularly.**

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

0.2 Table of Contents

- Section **??**
- Section **??**
- Section **??**
- Section **??**

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the RUBRIC.

Part I - Probability

To get started, let's import our libraries.

```
In [1]: import pandas as pd
        import numpy as np
        import random
        import matplotlib.pyplot as plt
        %matplotlib inline
        #We are setting the seed to assure you get the same answers on quizzes as we set up
        random.seed(42)
        import warnings
        warnings.filterwarnings('ignore')
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. **Use your dataframe to answer the questions in Quiz 1 of the classroom.**

a. Read in the dataset and take a look at the top few rows here:

```
In [2]: df = pd.read_csv('ab_data.csv')
        df.head(55)
```

Out[2]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
5	936923	2017-01-10 15:20:49.083499	control	old_page	0
6	679687	2017-01-19 03:26:46.940749	treatment	new_page	1
7	719014	2017-01-17 01:48:29.539573	control	old_page	0
8	817355	2017-01-04 17:58:08.979471	treatment	new_page	1
9	839785	2017-01-15 18:11:06.610965	treatment	new_page	1
10	929503	2017-01-18 05:37:11.527370	treatment	new_page	0
11	834487	2017-01-21 22:37:47.774891	treatment	new_page	0
12	803683	2017-01-09 06:05:16.222706	treatment	new_page	0
13	944475	2017-01-22 01:31:09.573836	treatment	new_page	0
14	718956	2017-01-22 11:45:11.327945	treatment	new_page	0
15	644214	2017-01-22 02:05:21.719434	control	old_page	1
16	847721	2017-01-17 14:01:00.090575	control	old_page	0
17	888545	2017-01-08 06:37:26.332945	treatment	new_page	1
18	650559	2017-01-24 11:55:51.084801	control	old_page	0
19	935734	2017-01-17 20:33:37.428378	control	old_page	0
20	740805	2017-01-12 18:59:45.453277	treatment	new_page	0
21	759875	2017-01-09 16:11:58.806110	treatment	new_page	0
22	767017	2017-01-12 22:58:14.991443	control	new_page	0
23	793849	2017-01-23 22:36:10.742811	treatment	new_page	0
24	905617	2017-01-20 14:12:19.345499	treatment	new_page	0
25	746742	2017-01-23 11:38:29.592148	control	old_page	0
26	892356	2017-01-05 09:35:14.904865	treatment	new_page	1
27	773302	2017-01-12 08:29:49.810594	treatment	new_page	0
28	913579	2017-01-24 09:11:39.164256	control	old_page	1
29	736159	2017-01-06 01:50:21.318242	treatment	new_page	0
30	690284	2017-01-13 17:22:57.182769	control	old_page	0
31	826115	2017-01-05 11:27:16.756633	treatment	new_page	0
32	875124	2017-01-05 15:39:25.439906	treatment	new_page	1
33	931013	2017-01-07 03:23:57.932344	treatment	new_page	0
34	710349	2017-01-11 22:24:44.226492	control	old_page	0
35	677533	2017-01-23 17:48:50.491821	control	old_page	0
36	831737	2017-01-11 21:18:20.911015	control	old_page	1
37	648583	2017-01-19 09:03:05.545308	treatment	new_page	0
38	728086	2017-01-03 17:07:00.837852	treatment	new_page	0

39	870163	2017-01-02 21:33:49.325594	treatment	new_page	0
40	771087	2017-01-16 00:05:29.983919	control	old_page	0
41	739414	2017-01-03 13:25:55.139705	treatment	new_page	0
42	896163	2017-01-22 09:10:20.753218	control	old_page	0
43	862225	2017-01-08 14:49:37.335432	control	old_page	1
44	939593	2017-01-05 09:15:31.984283	control	old_page	0
45	702260	2017-01-18 13:55:31.488221	control	old_page	0
46	943635	2017-01-22 13:37:39.722775	treatment	new_page	0
47	800436	2017-01-20 07:47:47.224386	treatment	new_page	0
48	698590	2017-01-23 11:51:59.925413	treatment	new_page	0
49	830513	2017-01-12 00:50:01.470557	treatment	new_page	0
50	670941	2017-01-05 08:16:41.306478	control	old_page	0
51	850231	2017-01-18 17:18:04.790584	control	old_page	1
52	916511	2017-01-22 06:20:04.691382	treatment	new_page	0
53	897174	2017-01-17 02:03:25.962173	treatment	new_page	0
54	906999	2017-01-22 19:16:44.715266	treatment	new_page	0

b. Use the cell below to find the number of rows in the dataset.

```
In [3]: len(df.index)
```

```
Out[3]: 294478
```

```
In [4]: df.user_id.nunique()
```

```
Out[4]: 290584
```

c. The number of unique users in the dataset.

```
In [12]: df['user_id'].nunique()
```

```
Out[12]: 290584
```

d. The proportion of users converted.

```
In [5]: df.query('converted == 1').user_id.nunique() / df.user_id.nunique()
```

```
Out[5]: 0.12104245244060237
```

e. The number of times the `new_page` and `treatment` don't match.

```
In [6]: df.query("(group != 'treatment' and landing_page == 'new_page') or (group == 'treatment'
```

```
Out[6]: 3893
```

f. Do any of the rows have missing values?

```
In [7]: df.shape[0] - df.dropna().shape[0]
```

```
Out[7]: 0
```

2. For the rows where **treatment** does not match with **new_page** or **control** does not match with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to figure out how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
In [ ]:
```

```
In [44]: df2 = df.query("(group == 'control' and landing_page == 'old_page') or (group == 'treat
         len(df2)
```

```
Out[44]: 290585
```

3. Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_id**s are in **df2**?

```
In [45]: df2.user_id.nunique()
```

```
Out[45]: 290584
```

b. There is one **user_id** repeated in **df2**. What is it?

```
In [10]: print(df2[df2['user_id'].duplicated()]['user_id'].to_string(index=False))
```

```
773192
```

c. What is the row information for the repeat **user_id**?

```
In [11]: df2[df2.duplicated(['user_id'])]
```

```
Out[11]:       user_id                   timestamp      group landing_page  converted
         2893   773192  2017-01-14 02:55:59.590927  treatment     new_page          0
```

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```
In [12]: df2.drop_duplicates(['user_id'], keep='first',inplace=True)
```

4. Use **df2** in the cells below to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
In [13]: df2['converted'].mean()
```

```
Out[13]: 0.11959708724499628
```

b. Given that an individual was in the `control` group, what is the probability they converted?

```
In [14]: df2[df2['group'] == 'control']['converted'].mean()
```

```
Out[14]: 0.1203863045004612
```

c. Given that an individual was in the `treatment` group, what is the probability they converted?

```
In [15]: df2[df2['group'] == 'treatment']['converted'].mean()
```

```
Out[15]: 0.11880806551510564
```

d. What is the probability that an individual received the new page?

```
In [16]: len(df2[df2['landing_page'] == 'new_page'].index)/len(df2.index)
```

```
Out[16]: 0.5000619442226688
```

e. Consider your results from parts (a) through (d) above, and explain below whether you think there is sufficient evidence to conclude that the new treatment page leads to more conversions.

Your answer goes here.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new}, which are the converted rates for the old and new pages.

Put your answer here.

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **conversion rate** for p_{new} under the null?

```
In [17]: p_new = df2['converted'].mean()
         p_new
```

```
Out[17]: 0.11959708724499628
```

b. What is the **conversion rate** for p_{old} under the null?

```
In [18]: p_old = df2['converted'].mean()
         p_old

Out[18]: 0.11959708724499628
```

c. What is n_{new}, the number of individuals in the treatment group?

```
In [19]: n_new = len(df2.query('landing_page == "new_page"'))
         n_new

Out[19]: 145310
```

d. What is n_{old}, the number of individuals in the control group?

```
In [20]: n_old = len(df2.query('landing_page == "old_page"'))
         n_old

Out[20]: 145274
```

e. Simulate n_{new} transactions with a conversion rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
In [21]: new_page_converted = np.random.binomial(1, p=p_new, size=n_new)
```

f. Simulate n_{old} transactions with a conversion rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
In [22]: old_page_converted = np.random.binomial(1, p=p_old, size=n_old)
```

g. Find p_{new} - p_{old} for your simulated values from part (e) and (f).

```
In [23]: new_page_converted.mean()-old_page_converted.mean()

Out[23]: -0.00094507911654652388
```

h. Create 10,000 p_{new} - p_{old} values using the same simulation process you used in parts (a) through (g) above. Store all 10,000 values in a NumPy array called **p_diffs**.

```
In [24]: new_converted_simulation = np.random.binomial(n_new, p_new,  10000)/n_new
         old_converted_simulation = np.random.binomial(n_old, p_old,  10000)/n_old
         p_diffs = new_converted_simulation - old_converted_simulation
```

i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
In [25]: plt.hist(p_diffs);
```

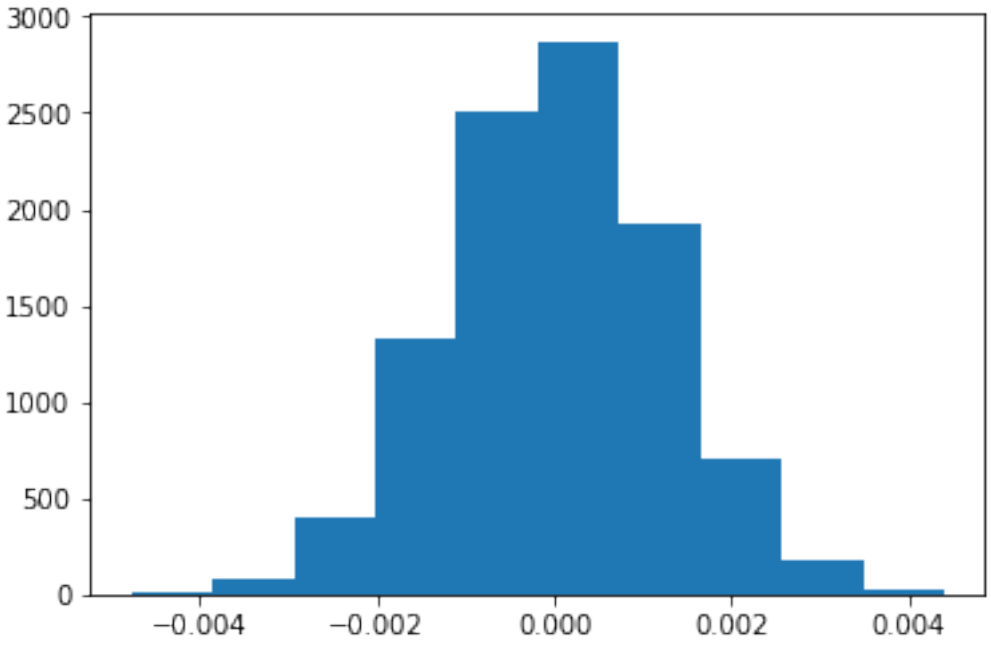

j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

```
In [26]: treatment_converted = df2.query('group =="treatment"').converted.mean()
         control_converted = df2.query('group =="control"').converted.mean()
         act_diff = treatment_converted - control_converted

         p_diffs = np.array(p_diffs)

         (p_diffs > act_diff).mean()
```

```
Out[26]: 0.90400000000000003
```

k. In words, explain what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

Answer: We calculated the p value. P value helps us determine the significance of our results, whether we accept the null hypothesis or reject it. Since the value p is not small, we could not reject the null hypothesis and do not prefer the new page and keep the old page.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let n_old and n_new refer the the number of rows associated with the old page and new pages, respectively.

len(df2.query(" landing_page == 'old_page' and converted == 1").index)

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the the number of rows associated with the old page and new pages, respectively.

```
In [28]: import statsmodels.api as sm

         convert_old = len(df2.query(" landing_page == 'old_page' and converted == 1").index)
         convert_new = len(df2.query(" landing_page == 'new_page' and converted == 1").index)
         n_old = len(df2[df2['group'] == 'control'].index)
         n_new = len(df2[df2['group'] == 'treatment'].index)
```

m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. Here is a helpful link on using the built in.

```
In [29]: z_score, p_value = sm.stats.proportions_ztest([convert_old, convert_new], [n_old, n_new

         z_score, p_value

Out[29]: (1.3109241984234394, 0.90505831275902449)
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts **j.** and **k.**?

```
In [30]: from scipy.stats import norm
         print(norm.cdf(z_score))
         print(norm.ppf(1-(0.05)))

0.905058312759
1.64485362695
```

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Logistic Regression

b. The goal is to use **statsmodels** to fit the regression model you specified in part **a.** to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create in df2 a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

```
In [54]: df2['intercept']=1
         df2[['control', 'ab_page']]=pd.get_dummies(df2['group'])
         df2.drop(labels=['control'], axis=1, inplace=True)
         df2.head()
```

Out[54]:

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	ab_page
0	1	0
1	1	0
2	1	1
3	1	1
4	1	0

c. Use **statsmodels** to instantiate your regression model on the two columns you created in part b., then fit the model using the two columns you created in part **b.** to predict whether or not an individual converts.

```
In [55]: import statsmodels.api as sm
         m = sm.Logit(df2.converted, df2[['intercept', 'ab_page']])
         results =  m.fit()
```

```
Optimization terminated successfully.
         Current function value: 0.366118
         Iterations 6
```

```
In [60]: from scipy import stats
         stats.chisqprob = lambda chisq, df: stats.chi2.sf(chisq, df)
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
In [61]: results.summary()
```

Out[61]:

```
<class 'statsmodels.iolib.summary.Summary'>
"""
                           Logit Regression Results
==============================================================================
Dep. Variable:              converted   No. Observations:               290585
Model:                          Logit   Df Residuals:                   290583
Method:                           MLE   Df Model:                            1
Date:                Mon, 28 Sep 2020   Pseudo R-squ.:               8.085e-06
Time:                        04:43:21   Log-Likelihood:            -1.0639e+05
```

```
converged:                       True   LL-Null:                       -1.0639e+05
                                        LLR p-value:                       0.1897
==============================================================================
                 coef    std err          z      P>|z|      [0.025      0.975]
------------------------------------------------------------------------------
intercept     -1.9888      0.008   -246.669      0.000      -2.005      -1.973
ab_page       -0.0150      0.011     -1.312      0.190      -0.037       0.007
==============================================================================
"""
```

e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint**: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in **Part II**?

**The p-value associated with ab_page is 0.19. It is not statistifically significant as it is not less than 0.05. The reason why p-value is different in Part 2 because in the previous part we performed a one-sided test, where in this part, logistic regression, we performed two-sided test. This implies that p_new is equal to p_old. which is the null hypothesis of a two tailed test.
Ho: p_new = p_old
H1: p_new != p_old**

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

Considering other things that might influence the conversions and taking them into account is a good idea. These new features can contribute to the significance of the results of our tests and lead to more precise decisions. One of the drawbacks of adding additional terms in the regression model is Simpson's paradox. It is that a trend can appear in several different groups of data but disappears or reverses when these groups are combined.

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in. You will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows. Here are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables.** Provide the statistical output as well as a written response to answer this question.

```
In [62]: countries_df = pd.read_csv('./countries.csv')
         df_new = countries_df.set_index('user_id').join(df2.set_index('user_id'), how='inner')
```

h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

```
In [63]: df_new[['UK', 'US']] = pd.get_dummies(df_new['country'])[['UK','US']]
         df_new.head()
```

Out[63]:

user_id	country	timestamp	group	landing_page	converted	intercept	ab_page	UK	US
630000	US	2017-01-19 06:26:06.548941	treatment	new_page	0	1	1	0	1
630001	US	2017-01-16 03:16:42.560309	treatment	new_page	1	1	1	0	1
630002	US	2017-01-19 19:20:56.438330	control	old_page	0	1	0	0	1
630003	US	2017-01-12 10:09:31.510471	treatment	new_page	0	1	1	0	1
630004	US	2017-01-18 20:23:58.824994	treatment	new_page	0	1	1	0	1

```
In [64]: model = sm.Logit(df_new.converted, df_new[['UK', 'US']])
         results = model.fit()
         results.summary()
```

```
Optimization terminated successfully.
         Current function value: 0.382863
         Iterations 6
```

Out[64]:

```
<class 'statsmodels.iolib.summary.Summary'>
"""
                           Logit Regression Results
==============================================================================
Dep. Variable:              converted   No. Observations:               290585
Model:                          Logit   Df Residuals:                   290583
Method:                           MLE   Df Model:                            1
Date:                Mon, 28 Sep 2020   Pseudo R-squ.:                -0.04573
Time:                        04:44:50   Log-Likelihood:            -1.1125e+05
converged:                       True   LL-Null:                   -1.0639e+05
                                        LLR p-value:                    1.000
==============================================================================
                 coef    std err          z      P>|z|      [0.025      0.975]
------------------------------------------------------------------------------
UK            -1.9868      0.011   -174.174      0.000      -2.009      -1.964
US            -1.9967      0.007   -292.315      0.000      -2.010      -1.983
==============================================================================
"""
```

Finishing Up

Congratulations! You have reached the end of the A/B Test Results project! You should be very proud of all you have accomplished!

Tip: Once you are satisfied with your work here, check over your report to make sure that it is satisfies all the areas of the rubric (found on the project submission page at the end of the lesson). You should also probably remove all of the "Tips" like this one so that the presentation is as polished as possible.

0.3 Directions to Submit

Before you submit your project, you need to create a .html or .pdf version of this notebook in the workspace here. To do that, run the code cell below. If it worked correctly, you should get a return code of 0, and you should see the generated .html file in the workspace directory (click on the orange Jupyter icon in the upper left).

Alternatively, you can download this report as .html via the **File** > **Download as** submenu, and then manually upload it into the workspace directory by clicking on the orange Jupyter icon in the upper left, then using the Upload button.

Once you've done this, you can submit your project by clicking on the "Submit Project" button in the lower right here. This will create and submit a zip file with this .ipynb doc and the .html or .pdf version you created. Congratulations!

In [ ]:
```
from subprocess import call
call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```